

Sofia Almeida

Berlin, Germany | sofia.almeida@example.com | Open to remote only

PROFILE

Lead data engineer with 11 years of experience designing high throughput ingestion platforms and mentoring engineering teams across distributed organisations. Has owned petabyte scale data infrastructure end to end, including cost, reliability and on call. Works remotely by preference and has led fully distributed teams across four time zones.

TECH STACK

Python, Scala, Spark, Kafka, Snowflake, dbt, Airflow, AWS, Terraform, Kubernetes, Flink, Databricks, Great Expectations, PostgreSQL

EXPERIENCE

Lead Data Engineer, Kestrel Systems, Berlin (2020 - Present)

- Rebuilt the company's batch ingestion platform on Spark and Kafka, raising daily throughput from 12 million to 90 million events while reducing infrastructure spend by 35 percent across three AWS regions.
- Migrated 240 legacy SQL transformations into dbt, cutting the nightly warehouse build from 7 hours to 95 minutes and giving analysts self service model documentation for the first time.
- Introduced Great Expectations data quality checks across 60 critical tables, reducing data incidents reported by downstream teams from roughly 9 per month to 2.
- Led a team of 6 engineers distributed across Berlin, Lisbon and Bangalore, running hiring, performance reviews and the on call rotation.
- Designed the streaming architecture that replaced a nightly batch with sub minute freshness for the fraud detection team, working with Flink and Kafka Streams.
- Reduced Snowflake spend by 28 percent through warehouse right sizing and clustering key changes, without any regression in query latency.

Senior Data Engineer, Alpenroute GmbH, Munich (2016 - 2020)

- Owned the Airflow deployment running 380 daily DAGs for the finance and logistics domains, including on call rotation and post incident review.
- Built a change data capture pipeline from 14 PostgreSQL databases into the warehouse, replacing nightly full table dumps that had grown to 6 hours.
- Mentored 4 junior engineers, two of whom were promoted to senior within two years.
- Established the code review and testing standards still used by the data team.

Data Engineer, Nordwind Logistics, Hamburg (2014 - 2016)

- Built the first ETL pipelines feeding the company's reporting warehouse.
- Automated shipment reconciliation across 3 carrier APIs.

LEADERSHIP AND COMMUNITY

Organiser, Berlin Data Engineering meetup (2019 - Present)

Speaker, Data Council Europe: Streaming Without the Tears (2023)

Mentor, Women in Data Berlin mentorship programme (2021 - Present)

EDUCATION

Diplom Informatik, TU Berlin (2012)