

# KUMAR ROHIT

Pune, Maharashtra • [kumarrohit2905@outlook.com](mailto:kumarrohit2905@outlook.com) • +91 6204767334 • [linkedin.com/in/kumarrohit25](https://linkedin.com/in/kumarrohit25) • [github.com/k-Rohit](https://github.com/k-Rohit)

## SUMMARY

Data Engineer (1.5+ years) specializing in GenAI and Agentic AI built production Databricks pipelines processing 2M+ records and improving performance by up to 75%, plus RAG/agentic systems using Python, LangChain, LangGraph, and hybrid search.

## PROJECTS

### [Arxiv Paper Curator — Agentic RAG & Hybrid Search System](#)

Python, FastAPI, LangGraph, OpenSearch, OpenAI, Docker, Airflow, LangSmith

- Architected an agentic RAG system for arXiv papers, incorporating citation-backed answers and conversation threads with PostgreSQL checkpointing for enhanced data retrieval.
- Implemented a hybrid retrieval pipeline with Docling PDF parsing, OpenAI embeddings, OpenSearch BM25 + HNSW search, and RRF fusion.
- Added observability with LangSmith traces and golden-set evals covering routing, retrieval recall, and answer correctness.

### [AgentCare](#)

LangGraph, FastAPI, Supabase, LangChain

- Built an agentic AI assistant for hospital workflows to handle appointment booking, rescheduling, document filing, reminders, and department routing.
- Architected persistent conversation state solutions with role-based access control and audit logging, ensuring secure and efficient healthcare administration workflows.

## SKILLS

**AI/ML & GenAI:** LangChain, LangGraph, Agentic AI, RAG, Hybrid Search, Prompt Engineering, LLM Evals, LangSmith

**Data Engineering:** Databricks, PySpark, ETL, Unity Catalog

**Databases & Search:** OpenSearch, Pinecone, FAISS, MySQL, SQL Server, PostgreSQL

**Languages & Frameworks:** Python, SQL, FastAPI, RESTful APIs

**Cloud & Tools:** Azure, Docker, Git, GitHub Actions, CI/CD

## EXPERIENCE

### Data Engineer, Philips, Chennai

Nov'25 - Present

- Automated TrackWise data extraction and SQL transformations, reducing a 1-hour daily process to a 15-minute pipeline, achieving 75% efficiency.
- Refactored downstream queries and pipelines to absorb schema evolutions across 25 reporting pipelines with zero reporting disruption.
- Migrated a legacy R-based complaint analytics solution to Databricks, cutting processing time 75% (8h to 2h).
- Architected end-to-end Databricks pipelines for automated regulatory and quality reporting, supporting 5 downstream reports.
- Implemented CI/CD deployment frameworks using Declarative Asset Bundles and GitHub Actions, ensuring uniform promotion across Dev, QA, and Prod

### Data Engineer, StoneX Group Inc., Pune

Jan'25 – Oct'25

- Built scalable ETL pipelines on Databricks, ingesting 2M+ records from SQL Server into a Medallion Architecture (Bronze/Silver/Gold) using Unity Catalog.
- Migrated a critical Airflow DAG delivering company/account data to Salesforce CRM into a Databricks workflow, reducing latency by 40%.
- Eliminated daily pipeline failures and improved CRM data integration accuracy through the Databricks-based redesign. [bullet point here.](#)

## EDUCATION

### Symbiosis Institute of Technology

2021 – 2025

B.Tech Artificial Intelligence and Machine Learning